

OPERATING SYSTEMS (THEORY)

LECTURE - 7

K.ARIVUSELVAN

Assistant Professor (Senior) – (SITE)

VIT University

PROCESS SYNCHRONIZATION

Introduction to Cooperating Processes

- Processes within a system may be **independent or cooperating**

- Independent process **cannot affect** or be **affected by** the **execution** of another process

- Cooperating process **can affect** or be **affected by** other processes, including **sharing data**

Race Condition:

- ❑ Several processes **access and manipulate** the **shared data** **concurrently**
- ❑ The **outcome of the execution** depends on the **particular order** in which the access takes place

time	Person A	Person B
8:00	Look in fridge. <i>Out of milk</i>	
8:05	Leave for store.	
8:10	Arrive at store.	Look in fridge. <i>Out of milk</i>
8:15	Buy milk.	Leave for store.
8:20	Leave the store.	Arrive at store.
8:25	Arrive home, put milk away.	Buy milk.
8:30		Leave the store.
8:35		Arrive home, <i>OH! OH!</i>

Someone gets milk, but NOT everyone (*too much milk!*)

Race Condition: Solution

☐ To prevent race conditions, **concurrent processes** must be **synchronized**

☐ Ensure that only **one process** at a time can be **manipulating the shared Data**

The Critical-Section

❑ A **section of code** Common to **n cooperating processes**, in which the processes may be accessing **common variables**

A critical section environment contains:

1. **Entry Section:** Code **requesting entry** into the critical section

2. **Critical Section:** Code in which **only one process can execute at any one time**

3. **Exit Section:** The **end** of the **critical section**, releasing or allowing others in

4. **Remainder Section:** Rest of the code **AFTER** the critical section

General Structure of a Typical Process

```
do {  
    entry section  
  
    critical section  
  
    exit session  
  
    remainder section  
  
} while (TRUE);
```


Solution to Critical-Section Problem

The critical section must **ENFORCE** all the **3 rules**:

(1) Mutual
Exclusion

(2) Progress

(3) Bounded
Waiting

Solution to Critical-Section Problem

1. Mutual Exclusion – If process P_i is **executing in its critical section**, then **no other processes** can be **executing in their critical sections**.

(i.e. **no two processes** will **simultaneously** be inside the **same CS**)

2. Progress – the **selection of the process** that will enter the critical section next **cannot be postponed indefinitely**.

3. Bounded Waiting – A bound must exist on the **number of times** that **other processes are allowed** to enter their **critical sections**.

+

➤ processes will **remain inside its CS** for a **short time only**, without **blocking**

Initial Attempts to Solve Problem

- ❑ Only 2 processes, P_0 and P_1

General structure of process P_i (other process P_j)

```
do {  
    entry section  
        critical section  
    exit section  
        remainder section  
} while (1);
```

- ❑ Processes may share some common variables to synchronize their actions.

Algorithm 1

□ Shared variables:

```
int turn;  
initially turn = 0 (or ) 1
```

□ Structure of Process P_i :

```
turn = j
```

```
( $P_j$  can enter its critical section)
```

```
repeat
```

```
    while (turn != i) do nothing; /*busy wait*/
```

```
    critical section
```

```
    turn = i;
```

```
    remainder section
```

```
Until false;
```

Algorithm 1

□ This solution guarantees **mutual exclusion**

Drawback 1: processes must **strictly alternate**

Drawback 2: if one processes fails other process is **permanently blocked**

Algorithm 2

Shared variables:

```
boolean flag[2];  
initially flag [0] = flag [1] = false
```

$\text{flag}[i] = \text{true} \Rightarrow P_i$ ready to enter its critical section

Structure of Process P_i :

```
repeat  
    flag[ i ] := true;  
    while (flag[ j ]) do nothing;  
    critical section  
    flag [j] = false;  
    remainder section  
Until false;
```

Algorithm 2

- ❑ This solution guarantees **mutual exclusion**

Drawback 1: This approach may lead to **dead lock**

- ❑ What is **wrong** with this implementation ?

➤ Dead lock occurs because each process can **insist on its right** to enter critical section

Algorithm 3 (PETERSON ALGORITHM)

❑ **Combined shared variables** of algorithms 1 and 2

Process P_i

```
do {  
    flag [i] := true;  
    turn = j;  
    while (flag [j] and turn = j) do nothing;  
    critical section  
    flag [i] = false;  
    remainder section  
} while (1);
```

❑ **Meets all three requirements;** solves the **critical-section problem** for two processes

Bakery Algorithm

Critical section for n processes:

- ❑ Before entering its critical section, process **receives a number**.
- ❑ Holder of the **smallest number** enters the **critical section**.
- ❑ If processes P_i and P_j receive the **same number**, if $i < j$, then P_i is **served first**; else P_j is served first.
- ❑ The **numbering scheme** always generates numbers in **increasing order** of enumeration; i.e., 1,2,3,3,4,5...

Notation => lexicographical order (**ticket #**, **process id #**)

□ **Shared data:**

```
boolean choosing[n];
```

```
int number[n];
```

- Data structures are initialized to **false** and **0** respectively

- **choosing** array => indicate that a process wants to enter it's critical section

- **number** array => contains the numbers associated with each process.

Bakery Algorithm

Repeat {

```
choosing[i] = true;  
number[i] = max(number[0], number[1], ..., number [n - 1])+1;  
choosing[i] = false;
```

```
for (j = 0 to n-1) do begin  
{  
while (choosing[j]) do no-op ;  
while ((number[j] != 0) && ((number[ j ], j ) < (number[ i ], i ) ))  
do no-op ;  
}
```

critical section

```
number[i] = 0;
```

remainder section

} until false;

Semaphores

□ Semaphore is a **variable** that has an **integer value**

➤ May be initialized to a **nonnegative number**

□ Can only be accessed via two indivisible (atomic) operations

wait (S):

S --;

signal (S):

S ++;

➤ **Wait** operation **decrements** the semaphore value

➤ **Signal** operation **increments** semaphore value

Semaphores

□ If a process is **waiting for a signal**, it is **suspended** until that signal is sent

□ **Queue** is used to **hold processes waiting** on the semaphore

Critical Section of N Processes

Shared data:

```
semaphore mutex; //initially mutex = 1
```

Process P_i :

```
do {  
    wait(mutex);  
  
    critical section  
  
    signal(mutex);  
  
    remainder section  
  
} while (1);
```

Implementation

□ Semaphore operations are now defined as:

wait(mutex):

```
mutex.value--;  
if (mutex.value < 0)  
{  
    add this process to S.L;  
    block;  
}
```

Implementation

signal(mutex):

```
mutex.value++;  
if (mutex.value <= 0)  
{  
    remove a process P from S.L;  
    wakeup(P);  
}
```

□ Value of semaphore can be **negative** and represents the **number of processes waiting** on it

Semaphore As a General Synchronization Tool

□ Execute **B** in **P_j** only after **A** executed in **P_i**

□ Use semaphore *flag* initialized to 0

▪ Code:

P_i
:
A
signal(flag)


P_j
:
wait(flag)
B

Semaphores can be used to control the order in which concurrent processes run. In this example the following conditions must be met:-

- *Process A must run first
- *Process D can not run until after Process C is finished
- *Process E must not run until all other processes are finished

Note that semA is a general semaphore that is used to control the start of both Processes B and C.

LEGEND



Two Types of Semaphores

❑ **Counting semaphore** – integer value can range over an unrestricted domain

❑ **Binary semaphore** – integer value can range only between 0 and 1; can be simpler to implement

Classical Problems of Synchronization

☐ **Bounded-Buffer Problem**

☐ **Dining-Philosophers Problem**

☐ **Readers-Writers Problem**

(1) Bounded-buffer Problem

Shared data:

semaphore **full**, **empty**, **mutex** ;

Initially:

full = 0, **empty** = n, **mutex** = 1

- **Buffer size** is n
- **Mutex** provides **exclusive access** to the **buffer**
- **Consumers** wait on **full**
- **Producers** wait on **empty**

Bounded-buffer Problem: **Producer Process**

do {

```
...  
produce an item in nextp  
...  
wait(empty);  
wait(mutex);  
...  
add nextp to buffer  
...  
signal(mutex);  
signal(full);
```

} while (1);

Bounded-buffer Problem: Consumer Process

do {

wait(full)
wait(mutex);

...
remove an item from **buffer** to **nextc**
...

signal(mutex);
signal(empty);

...
consume the item in **nextc**
...

} while (1);

(2) Dining-philosophers Problem



❑ To **start eating**, a philosopher needs **two chopsticks**

❑ **After eating**, the philosopher **releases** both the chopsticks

Shared data:

semaphore chopstick[5];
(Initially all values are 1)

Philosopher i :

do {

wait(chopstick[i])
wait(chopstick[(i+1) % 5])

...

eat

...

signal(chopstick[i]);
signal(chopstick[(i+1) % 5]);

...

think

...

} while (1);

❑ The solution is **not deadlock free!!**

=> All philosophers **pick up left chopsticks!!**

Solution:

- Allow **at most 4 philosophers** to be sitting on the table
- Allow a philosopher to **pick chopsticks only if both are available**
- An **odd philosopher** picks up **first the left** and then the **right chopstick** while a **even philosopher** does the **reverse**

(3) Readers-Writers Problem

❑ Two readers can access the shared data item simultaneously

❑ A writer requires exclusive access

❑ No reader should wait unless a writer is already in critical section

Shared data:

```
var mutex, wrt : semaphore (=1);  
    readcount : integer (=0);
```

Readers-Writers Problem

- ❑ *wrt* is common to both readers and writers.
 - ❑ It functions as mutual exclusion semaphore for writers.
 - ❑ It is also used by first and last reader that enters or exits the CS.
-
- ❑ *readcount* keeps track of how many readers are currently accessing the object.
-
- ❑ *mutex* provides mutual exclusion for updating *readcount*.

Readers-Writers Problem

Writer process

wait(wrt);

...

writing is performed

...

signal(wrt);

Readers-Writers Problem

Reader process

```
wait(mutex);  
    readcount := readcount + 1;  
    if readcount = 1 then wait(wrt);  
signal(mutex);  
    ...  
    reading is performed  
    ...  
wait(mutex);  
    readcount := readcount - 1;  
    if readcount = 0 then signal(wrt);  
signal(mutex);
```

Monitors

□ High-level **synchronization construct** that **allows the safe sharing** of an **abstract data type** among concurrent processes.

```
type monitor-name = monitor
    variable declarations
    procedure entry P1 :(...);
        begin ... end;
    procedure entry P2(...);
        begin ... end;
        ⋮
    procedure entry Pn (...);
        begin...end;
    begin
        initialization code
    end
```

Monitors

A monitor is similar to a C++ class that ties the data, operations, and in particular, the synchronization operations all together,

Unlike classes,

monitors guarantee mutual exclusion, i.e., only one thread may execute a given monitor method at a time.


Monitors

A Monitor defines a *lock* and zero or more *condition variables* for managing concurrent access to shared data.

The monitor uses the *lock* to insure that only a single thread is active in the monitor at any instance.

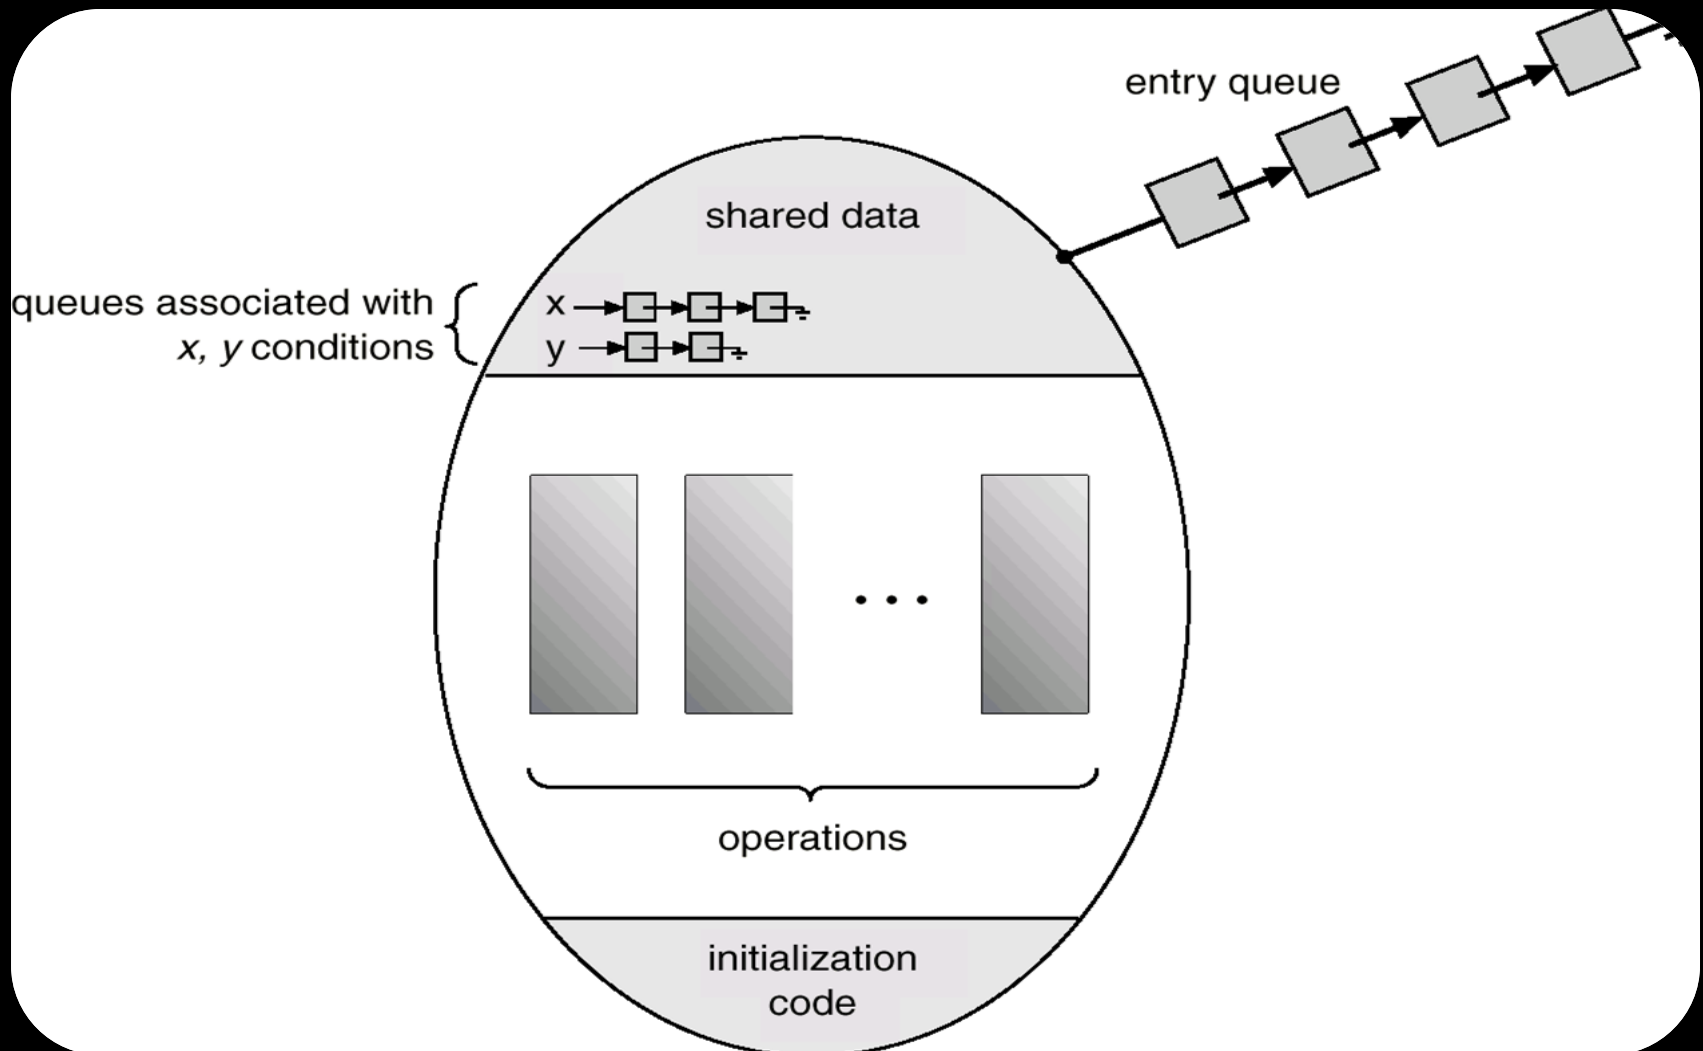
Condition variables enable threads to go to sleep inside of critical sections, by releasing their lock at the same time it puts the thread to sleep.

Schematic view of a monitor

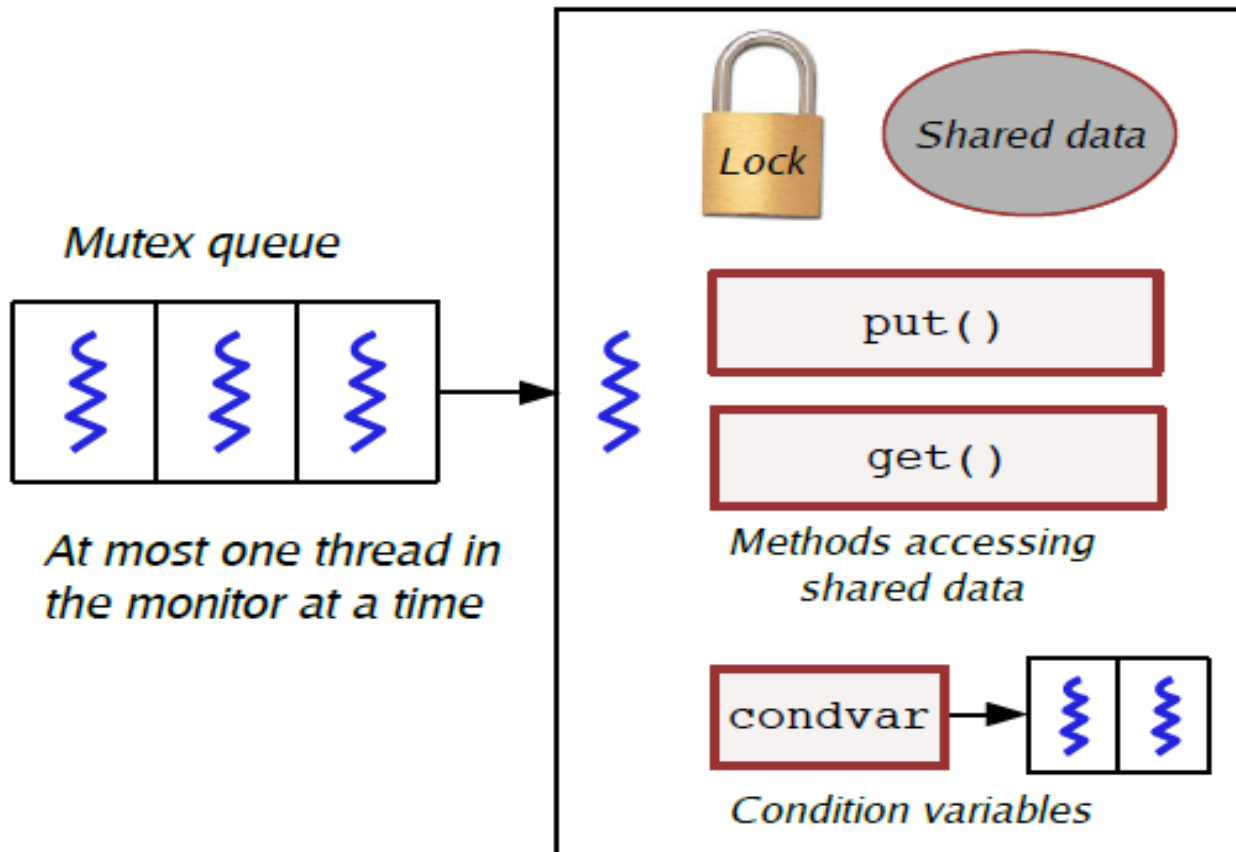


code
initialization

Monitor with condition variables



Monitor Example



Operations on Condition Variables

Condition variable: is a queue of threads waiting for something inside a critical section.


Condition variables support three operations:

Wait(Lock lock): atomic (release lock, go to sleep), when the process wakes up it re-acquires lock.

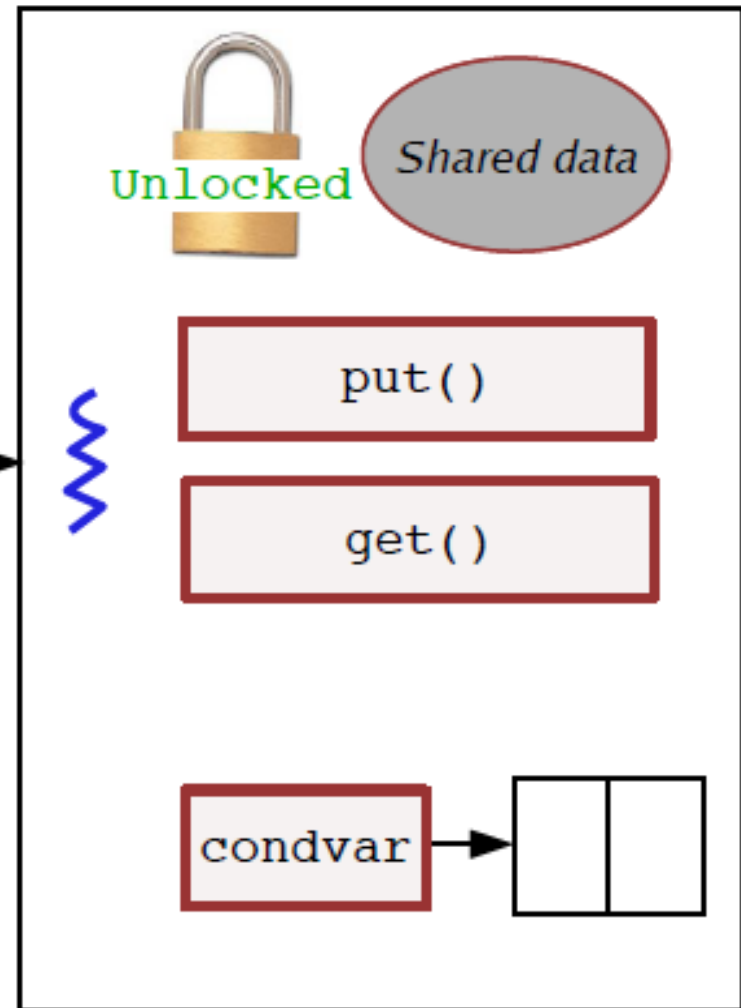
Signal(): wake up waiting thread, if one exists. Otherwise, it does nothing.

Broadcast(): wake up all waiting threads

Monitors



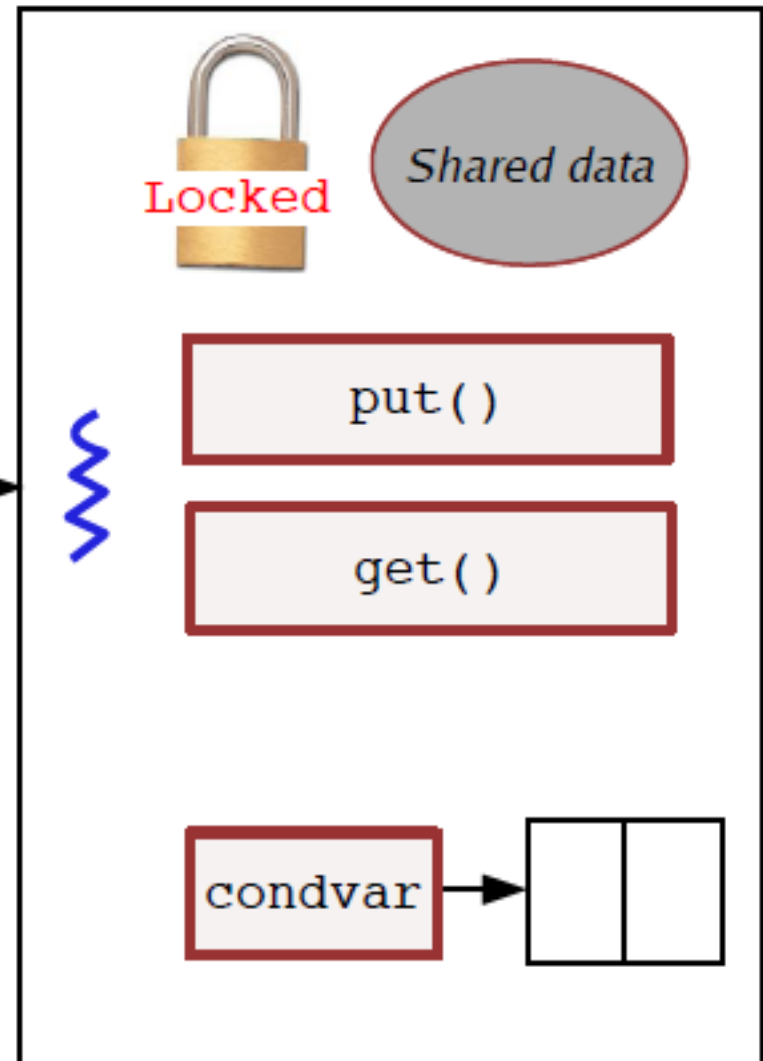
1) Thread enters monitor
(grabs lock)



Monitors



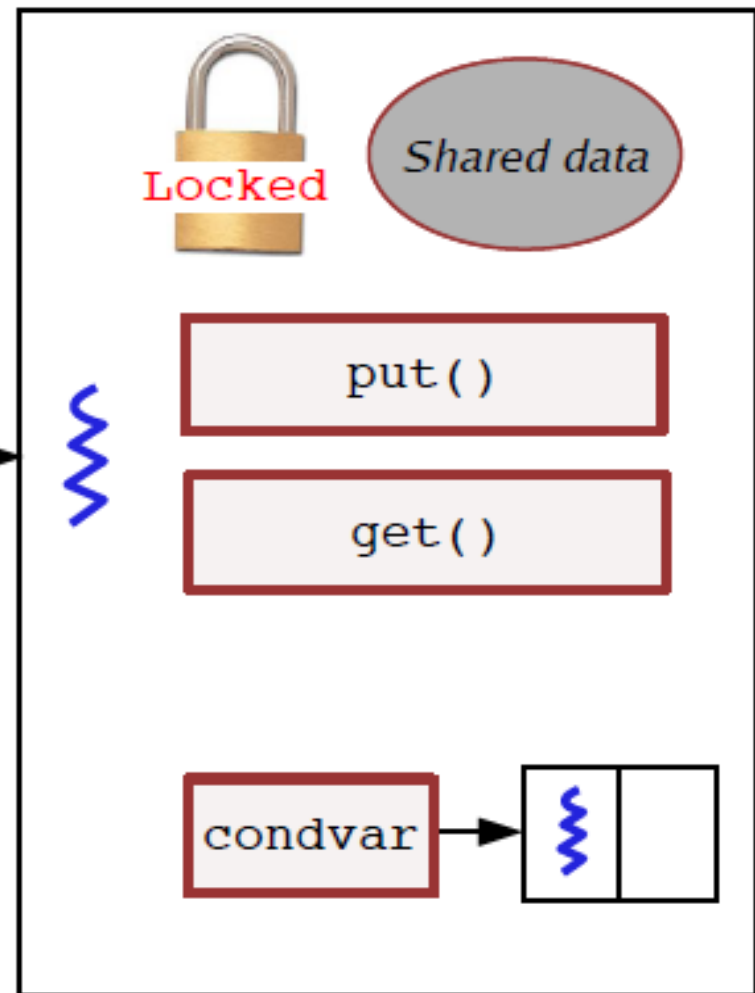
- 1) Thread enters monitor (grabs lock)
- 2) Other threads queue up



Monitors



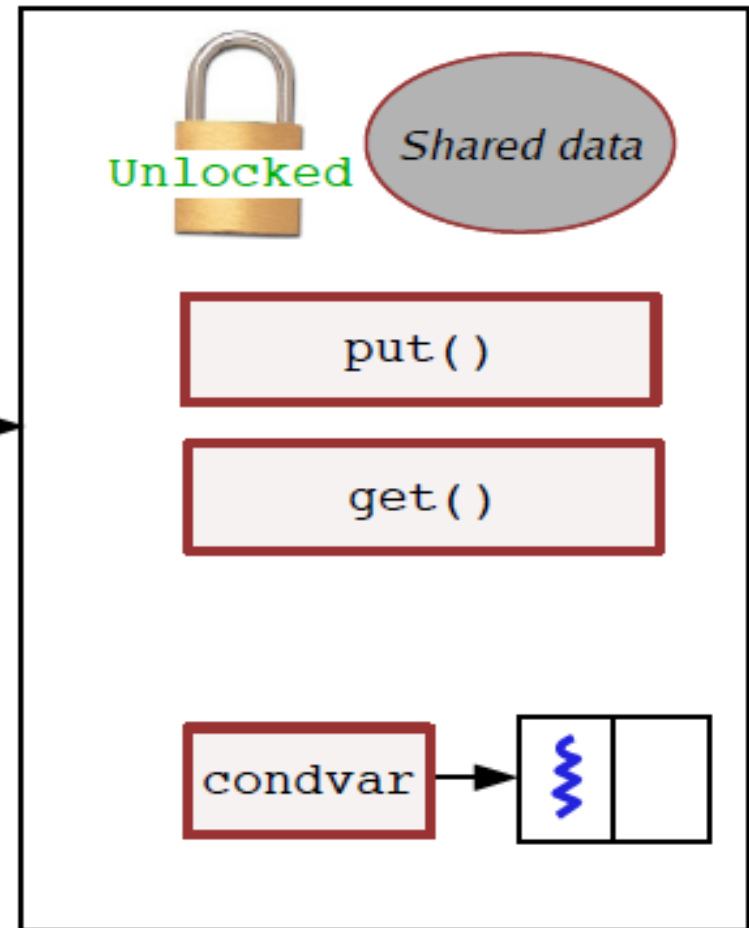
- 1) Thread enters monitor (grabs lock)
- 2) Other threads queue up
- 3) Blue thread waits() on CV



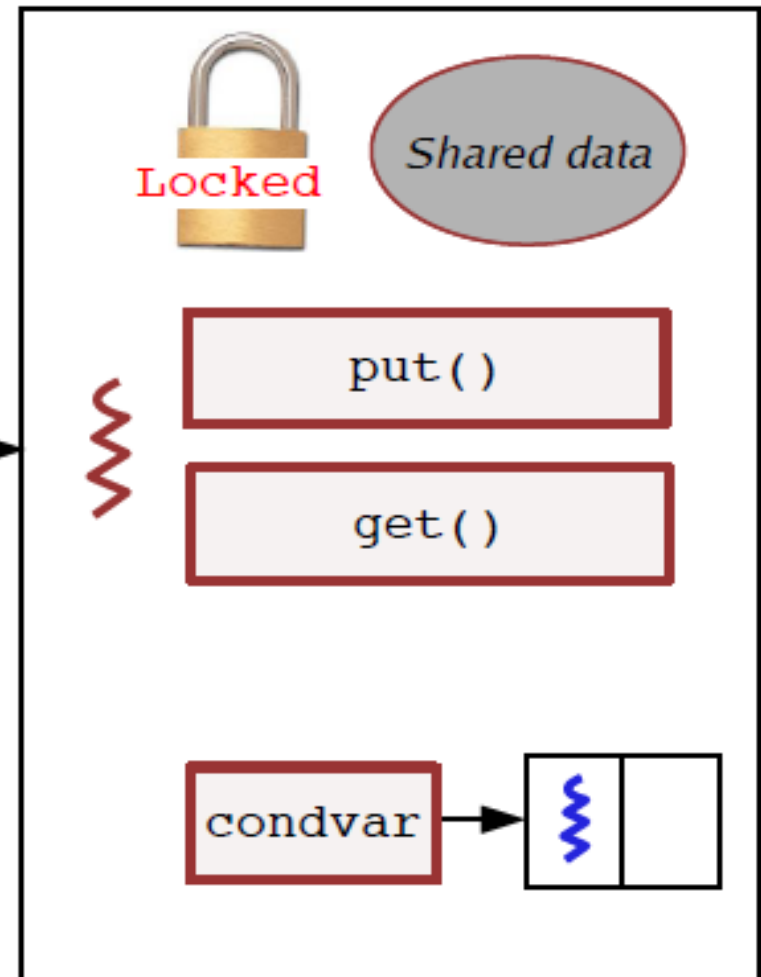

Monitors



- 1) Thread enters monitor (grabs lock)
- 2) Other threads queue up
- 3) Blue thread waits() on CV
- 4) Next thread enters monitor



Monitors

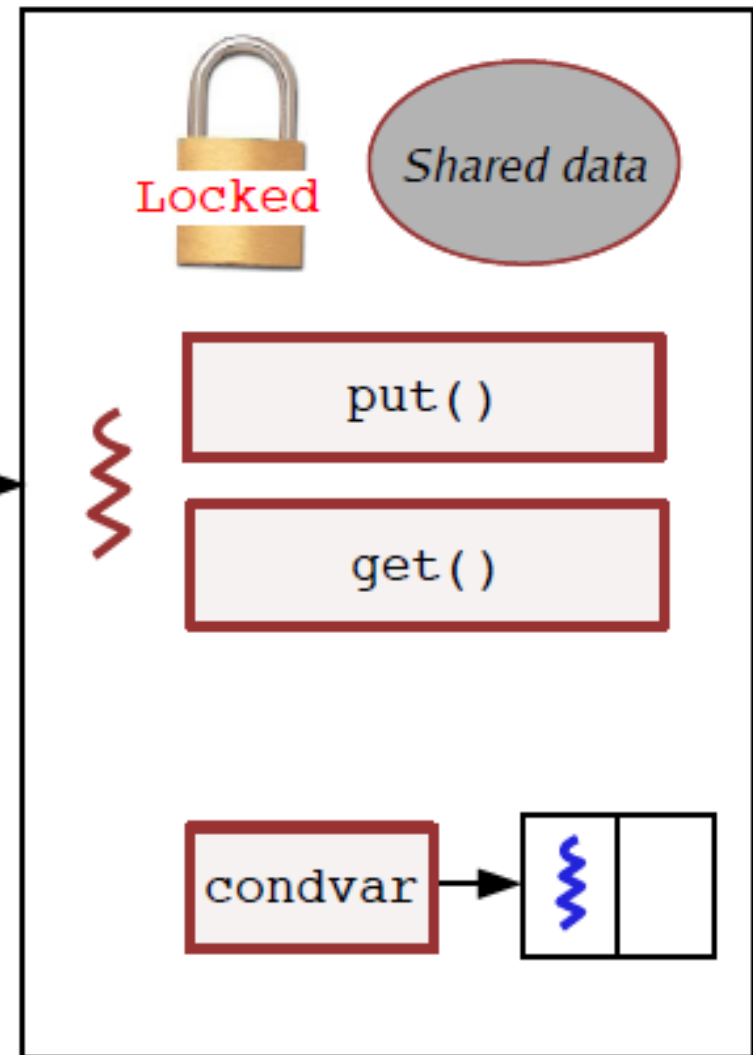


- 1) Thread enters monitor (grabs lock)
- 2) Other threads queue up
- 3) Blue thread waits() on CV
- 4) Next thread enters monitor


Monitors



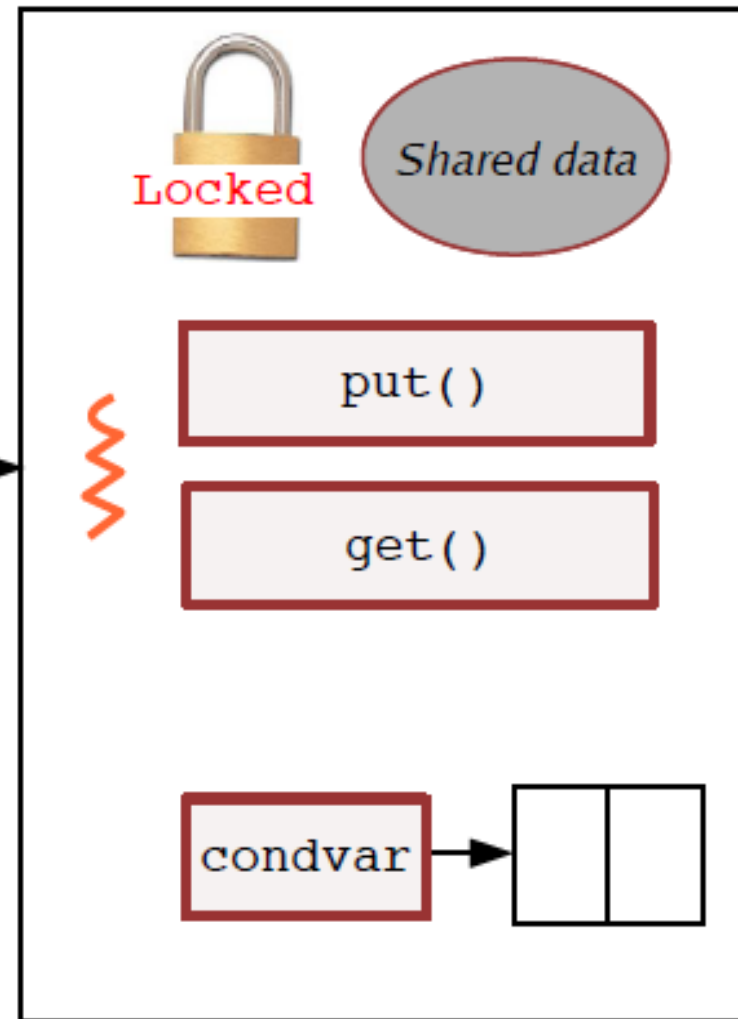
5) Thread in monitor calls
notify() on CV



Monitors



- 5) Thread in monitor calls `notify()` on CV
- 6) Next thread enters monitor (order depends on lock implementation!)



Dining-philosophers Problem with Monitor

```
monitor dp
{
    enum {thinking, hungry, eating} state[5];
    condition self[5];
```

```
void pickup(int i)
void putdown(int i)
void test(int i)
```

```
void init() {
    for (int i = 0; i < 5; i++)
        state[i] = thinking;
}
```

Dining-philosophers Problem with Monitor

```
void pickup(int i) {  
    state[i] = hungry;  
    test[i];  
    if (state[i] != eating)  
        self[i].wait();  
}
```

```
void putdown(int i) {  
    state[i] = thinking;  
    /* test left and right neighbors */  
    test((i+4) mod 5);  
    test((i+1) mod 5);  
}
```

Dining-philosophers Problem with Monitor

```
void test(int i) {  
    if ( (state[(i + 4) mod 5] != eating) &&  
        (state[i] = hungry) &&  
        (state[(i + 1) mod 5] != eating)) {  
        state[i] = eating;  
        self[i].signal();  
    }  
}
```

Producer-Consumer Problem with Monitor

```
Monitor Producer-consumer
{
    condition full, empty;
    int count;
    void insert(int item); //the following slide
    int remove(); //the following slide
    void init() {
        count = 0;
    }
}
```


Producer-Consumer Problem with Monitor

```
void insert(int item)
{
    if (count == N)    full.wait();
        insert_item(item); // Add the new item
    count ++;
    if (count == 1)    empty.signal();
}
```

```
int remove()
{ int m;
  if (count == 0)    empty.wait();
  m = remove_item(); // Retrieve one item
  count --;
  if (count == N -1)    full.signal();
      return m;
}
```

Producer-Consumer Problem with Monitor

```
void producer()      //Producer process
{
    while (1) {
        item = Produce_Item();
        Producer-consumer.insert(item);
    }
}
```

```
void consumer()  //Consumer process
{
    while (1) {
        item = Producer-consumer.remove(item);
        consume_item(item);
    }
}
```

Monitor VS Semaphores

Semaphores

- ✧ wait()/signal() implement blocking mutual exclusion
- ✧ Also used as atomic counters (counting semaphores)
- ✧ Can be inconvenient to change and debug

Monitors

Synchronizes execution within procedures that manipulate encapsulated data shared among procedures

- ✧ Only one thread can execute within a monitor at a time
- ✧ Relies upon high-level language support

Monitor VS Semaphores

Semaphores (Disadvantages):

Low-level:

- Easy to forget a Set or a Reset of the Semaphore

Scattered:

- Semaphore calls are scattered in the code
- Difficult and error-prone to debug (aliasing!).