

OPERATING SYSTEMS (THEORY)

LECTURE - 8

K.ARIVUSELVAN

Assistant Professor (Senior) – (SITE)

VIT University

DEADLOCKS

The Deadlock Problem

□ Processes each **holding a resource** and **waiting to acquire a resource** held by another process

▪Example1

→ System has **2 tape drives**

→ P1 and P2 each **hold one tape drive** and each **needs another one**

Deadlock Characterization

Deadlock can arise if **four conditions hold** simultaneously:

(1) **Mutual exclusion:** only **one process at a time** can use a resource (i.e. resource is non-sharable).

(2) **Hold and wait:** a process **holding at least one** resource is **waiting to acquire additional resources** held by other processes.

Deadlock Characterization

(3) **No preemption:** a resource can be **released only voluntarily** by the process holding it, after that process has completed its task.

(4) **Circular wait:** there exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that:

- => P_0 is waiting for a resource that is held by P_1 ,
- => P_1 is waiting for a resource that is held by P_2, \dots ,
- => P_{n-1} is waiting for a resource that is held by P_n ,
- and
- => P_n is waiting for a resource that is held by P_0 .

Methods for Handling Deadlocks

➤ **Deadlock Prevention**

➤ **Deadlock Avoidance**

Deadlock Prevention

=> Ensure atleast one of the necessary conditions **cannot hold**.

Deadlock Avoidance

=> Given **additional information** in **advance**, concerning which **resources** a process will **request** and use during its **life time**

NOTE:

=> A System does **not employ** either a **deadlock prevention** or **deadlock avoidance** algorithm , it leads to a **deadlock situation**.

Deadlock Prevention

(1) Mutual Exclusion:

=> Hold for **non-sharable** resource (**Printer can't share simultaneously**)

=> **Sharable** resource mutual exclusion not required (**Read only file access by several process simultaneously**). No dead lock

=> In general it is **not possible** to **prevent deadlock** by **denying** the mutual exclusion condition.

Deadlock Prevention

(2) Hold & Wait:

1st Protocol :

=> Allot **all requested resource** before **begins execution**

Example:

Copying data from Tape drive to **Disk file**, Sort **Disk file**, Print

Drawback:

Hold printer **entire execution** even printer need **at last**

Deadlock Prevention

(2) Hold & Wait:

2nd Protocol :

=> Request **some resource** and **use** them

=> Before request, **release** additional resource currently allotted.

Example:

- **Request** Tape drive & Disk
- **Release**
- **Request** Disk & Printer

Deadlock Prevention

(3) No Preemption:

Protocol :

=> Check for **resource availability**

if (**available**)

{ **allocate** resource }

Else

{Check other process **hold resource** and **request additional resource** }

If (**Yes**) then **PREEMPT**

Deadlock Prevention

(4) Circular Wait:

Protocol :

=> Assign **unique integer** to **resource type**

Example:

Tape Drive = **1**, Disk = **4**, Printer = **10**

Constraint :


=> Process **request the resource** only in **increasing order**

=> Process that want to use **both Tape & Printer** at the **same time** must **first request** Tape & **then** Printer

Deadlock Avoidance

Safe State:

- When a process **requests** an **available resource**, system must **decide** if **immediate allocation** leaves the system in a **safe state**.
- System is in **safe state** if there exists a **safe sequence** of all processes.



Example – Safe State

▪ System with

12 Tape Drives &

3 Processes (P_0 , P_1 , P_2)

Time T_0 :

	Max.Needs	Allocation	Available
P_0	10	5	3
P_1	4	2	
P_2	9	2	

▪ Sequence $\langle P_1, P_0, P_2 \rangle$ is in Safe State or Unsafe State at Time T_0

Banker's Algorithm

- A resource allocation system with Multiple instances
- Each process must a priori claim maximum use.

Constraint:

- When a process requests a resource it may have to wait.
- When a process gets all its resources it must return them in a finite amount of time.

Data Structures for the Banker's Algorithm

Let n = number of processes, and m = number of resources types.

- **Available:** Number of available resources. If $\text{Available}[j] = k$, there are k instances of resource type R_j available.
- **Max:** Maximum demand of each process. If $\text{Max}[i,j] = k$, then process P_i may request at most k instances of resource type R_j .
- **Allocation:** Number of resource currently allocated to each process. If $\text{Allocation}[i,j] = k$ then P_i is currently allocated k instances of R_j .
- **Need:** Remaining resource need of each process. If $\text{Need}[i,j] = k$, then P_i may need k more instances of R_j to complete its task.

$$\text{Need}[i,j] = \text{Max}[i,j] - \text{Allocation}[i,j]$$

Example

- 5 processes **P0** through **P4**;
- 3 resource types **A** (10 instances), **B** (5 instances) and **C** (7 instances).

- Snapshot at time **T0**:

	Allocation	Max	Available
	A B C	A B C	A B C
P0	0 1 0	7 5 3	3 3 2
P1	2 0 0	3 2 2	
P2	3 0 2	9 0 2	
P3	2 1 1	2 2 2	
P4	0 0 2	4 3 3	

Example (Cont.)

- The content of the matrix **Need** is defined to be **Max – Allocation**.

	<u><i>Need</i></u>		
	<i>A</i>	<i>B</i>	<i>C</i>
<i>P0</i>	7	4	3
<i>P1</i>	1	2	2
<i>P2</i>	6	0	0
<i>P3</i>	0	1	1
<i>P4</i>	4	3	1

Resource-Request Algorithm for Process P_i

- **Request_i** = request vector for process P_i .
- If **Request_i [j] = k** then process P_i wants **k** instances of resource type R_j .

ACTION TAKEN:

1. If **Request_i ≤ Need_i** go to **step 2**. Otherwise, raise **error condition**, since process has **exceeded** its **maximum claim**.

2. If **Request_i ≤ Available**, go to **step 3**. Otherwise P_i must **wait**, since resources are **not available**.

3. Pretend to **allocate** requested resources to P_i by **modifying the state** as follows:

How?

Check if this **new state** is Safe.
i.e. if a safe sequence exists

$Available = Available - Request_i ;$

$Allocation_i = Allocation_i + Request_i ;$

$Need_i = Need_i - Request_i ;$

➤ If **safe** \Rightarrow the resources are **allocated** to P_i .

➤ If **unsafe** $\Rightarrow P_i$ must **wait**, and the **old resource-allocation state** is **restored**

Safety Algorithm

- Finding out whether or not a system is in safe state

1. Let **Work** and **Finish** be vectors of length **m** and **n**, respectively.
Initialize:

Work = **Available**

Finish [**i**] = **false** for **i** = 1,2,3, ..., **n**.

2. Find an **i** such that **both**:

(a) **Finish** [**i**] = **false**





(b) **Need**_{**i**} ≤ **Work**

If no such **i** exists, go to step 4.

3. **Work** = **Work** + **Allocation**_{**i**}
Finish[**i**] = **true**
go to step 2.b

4. If **Finish** [**i**] == **true** for all **i**, then the system is in a safe state.

Example (Cont.)

Seq	Avail	Need	New=Avail-Need	Max	Max + New
P ₁	3 3 2	1 2 2	2 1 0	3 2 2	5 3 2
P ₃					
P ₄					
P ₂					
P ₀					

Resource-allocation Graph

A **visual** (mathematical) way to determine if a deadlock has, or may occur.

A set of **vertices** V and a set of **edges** E

- V is partitioned into two types:

$\Rightarrow P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the **processes** in the system.


$\Rightarrow R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all **resource types** in the system

- **Request edge** – directed edge $P_i \rightarrow R_j$

- **Assignment edge** – directed edge $R_j \rightarrow P_i$

- If graph contains **cycles** \Rightarrow **deadlock**

Resource Allocation Graph



Wait-For Graph (Single Instance)


- Maintain *wait-for* graph

→ Nodes are processes

→ $P_i \rightarrow P_j$ if P_i is waiting for P_j

- Periodically invoke an algorithm that searches for a cycle in the graph.


Resource-Allocation Graph and Wait-for Graph



(a)

(g)

Resource-Allocation Graph



(b)

(p)

Corresponding wait-for graph

Deadlock Detection

- Allow system to **enter deadlock state**

(1) **Detection** algorithm

(2) **Recovery** scheme

Detection Algorithm

Several Instances of a Resource Type

Data Structures:

Available: Number of **available resources** of each type.

Allocation: The number of resources of each type **currently allocated** to each process.

Request: The **current request** of each process.

If **Request** $[i,j] = k$, then process P_i is requesting k more instances of resource type. R_j .

Detection Algorithm

1. Let *Work* and *Finish* be vectors of length *m* and *n*, respectively. Initialize:

(a) *Work* = *Available*

(b) For $i = 1, 2, \dots, n$,
if $\text{Allocation}_i \neq 0$, then $\text{Finish}[i] = \text{false}$;
otherwise, $\text{Finish}[i] = \text{true}$ (i.e., $\text{Allocation}_i = 0$).

2. Find an index *i* such that both:

(a) $\text{Finish}[i] == \text{false}$

(b) $\text{Request}_i \leq \text{Work}$

If no such *i* exists, go to step 4.

Detection Algorithm (Cont.)

3. $Work = Work + Allocation_i$
 $Finish[i] = true$
go to step 2.

4. If $Finish[i] == false$, for some i , $1 \leq i \leq n$, then the system is in deadlock state.

Moreover, if $Finish[i] == false$, then P_i is deadlocked.

Example of Detection Algorithm

- Five processes P_0 through P_4 ;
 - Three resource types **A** (7 instances), **B** (2 instances), and **C** (6 instances)
- Snapshot at time T_0 :

	<u>Allocation</u>	<u>Request</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	0 0 0	0 0 0
P_1	2 0 0	2 0 2	
P_2	3 0 3	0 0 0	
P_3	2 1 1	1 0 0	
P_4	0 0 2	0 0 2	

- Sequence $\langle P_0, P_2, P_3, P_1, P_4 \rangle$ will result in $Finish[i] = \text{true}$ for all i .

Example (Cont.)

- **P2** requests an additional instance of type **C**.

	<u>Request</u>
	A B C
P0	0 0 0
P1	2 0 1
P2	0 0 1
P3	1 0 0
P4	0 0 2

- **State of system?**

→ Can reclaim resources held by process **P0**, but insufficient resources to fulfill other processes requests.

→ **Deadlock exists**, consisting of processes **P1**, **P2**, **P3**, and **P4**.

Recovery from Deadlock: Process Termination

- Abort **all deadlocked processes**.
- Abort **one process at a time** until the deadlock cycle is eliminated.
- In **which order** should we choose to **abort**?

Priority of the process

How long process has **computed**, and **how much longer** to completion

Resources the process has **used**

Resources process **needs** to complete